

Parallel Programming with Python

Antonio Ruiz
Data Engineering
Universidad Politécnica de Yucatán
Ucú, Yucatán, México
2009121@upy.edu.mx

Abstract—This report outlines the numerical computation of π using different methods of parallel computing in Python. We explore the implementation and performance of each approach to understand the best practices in numerical integration techniques.

I. INTRODUCTION

The mathematical constant π (pi) is a fundamental element in various fields of science and engineering, including mathematics, physics, and engineering. It represents the ratio of the circumference of a circle to its diameter and appears in countless formulas across numerous disciplines. Despite its ubiquitous presence, π is an irrational number, meaning it cannot be expressed exactly as a finite decimal or a fraction, making its exact computation inherently challenging.

Numerical integration offers a practical method for approximating π by leveraging its geometric interpretation—the area of a circle. In particular, by approximating the area of a quarter circle inscribed within a unit square, one can derive an estimation for π . This method involves calculating the integral of the function $f(x) = \sqrt{1 - x^2}$ from 0 to 1, which geometrically represents a quarter circle. The exact area under this curve is $\pi/4$, and hence, numerical methods to approximate this integral can provide estimates for π .

This report explores the numerical approximation of π using Riemann sums, a basic approach in numerical integration. By subdividing the interval from 0 to 1 into smaller segments and summing the areas of rectangles under the curve, we can approximate the area of the quarter circle. The accuracy of this approximation depends significantly on the number of subdivisions or the width of each rectangle, denoted as Δx . As Δx decreases, or equivalently, as the number of intervals N increases, the approximation becomes more precise.

The simplicity of this method makes it ideal for implementation using various computational approaches, including sequential and parallel processing. This report not only implements these methods in Python but also evaluates their performance and efficiency, highlighting the advantages of parallel computing in handling computationally intensive tasks. The parallelization of this problem is particularly interesting due to its inherent divisibility and the independent nature of the calculations involved in each subdivision. This property allows for significant improvements in computation times, making it an excellent candidate for demonstrating the benefits of parallel and distributed computing.

A. Non-Parallel Solution

The first approach in our exploration of numerical methods for computing π involves a straightforward, non-parallel computation. This method is based on the classic Riemann sum approach to estimate the area under the curve of the function $f(x) = \sqrt{1 - x^2}$, which represents a quarter circle.

The implementation is encapsulated within a Python class, `PiCalculator`, which is designed to compute π by summing the areas of rectangles under the curve. The constructor of the class initializes the number of rectangles, N , which determines the fineness of the subdivision and, consequently, the accuracy of the approximation.

```
1 import numpy as np
2
3 class PiCalculator:
4     def __init__(self, N):
5         self.N = N # Number of rectangles
6
7     def compute_pi(self):
8         delta_x = 1 / self.N # Width of each small
9         rectangle
10        total_area = 0 # Accumulator for the sum of
11        areas of rectangles
12
13        # Loop to calculate area of each rectangle
14        for i in range(self.N):
15            x_i = i * delta_x # x-coordinate at the
16            left side of the rectangle
17            f_x = np.sqrt(1 - x_i**2) # Height of
18            the rectangle
19            total_area += delta_x * f_x # Add area
20            of the rectangle to total
21
22        # Multiply by 4 to get the area of the whole
23        circle
24        pi_approximation = 4 * total_area
25        return pi_approximation
26
27 # Example usage:
28 calculator = PiCalculator(N=100000)
29 pi_approx1 = calculator.compute_pi()
30 print(f"Approximation of \(\pi\) with N={calculator.
31 N}: {pi_approx1}")
```

Listing 1. Non-parallel Python Code for π Approximation

The key component of this implementation is the loop that iterates from 0 to $N - 1$, calculating the area of each rectangle. The area is computed by multiplying the width of the rectangle, Δx , by the height of the rectangle at the given x position, $f(x_i)$. This method, though simple and easy to understand, does not leverage any form of parallelization and thus serves as a baseline for evaluating more advanced techniques. The

performance and scalability of this approach are inherently limited by the sequential nature of the loop, making it a prime candidate for optimization through parallel computing techniques, which will be discussed in the following sections.

B. Results of the Non-Parallel Solution

After executing the non-parallel PiCalculator with $N = 100,000$, the program produced an approximation for π as follows:

Approximation of π with $N = 100,000$:
3.14161261640195

This result is quite close to the actual value of π ($\pi \approx 3.14159$), demonstrating the potential accuracy of the numerical integration method using Riemann sums. The approach, while simplistic and sequential in nature, is capable of providing a reasonable approximation of π with a sufficiently large number of subdivisions.

1) *Analysis:* The accuracy of this method is contingent upon the number of subdivisions N ; a higher N leads to a more precise approximation. The choice of $N = 100,000$ rectangles under the curve of $f(x) = \sqrt{1-x^2}$ balances computation time with precision. The error margin observed, when compared to the actual value of π , is approximately 0.00002, or 0.002%. This level of precision is adequate for many practical applications, though further reductions in error can be achieved by increasing N .

However, this increase in accuracy comes at the cost of increased computational resources and time, as each increase in N requires more iterations over the loop. This consideration is crucial in applications where computational resources or time are constrained.

The simplicity of the non-parallel approach serves as an excellent baseline for comparison with more complex parallel computing strategies, which are expected to reduce computation time significantly while maintaining or improving the accuracy of the approximation.

C. Parallel Computing with Multiprocessing

To enhance the computational efficiency of our π approximation, the second approach utilizes the multiprocessing module in Python. This method leverages multiple CPU cores to parallelize the computation of the areas of the rectangles used in the Riemann sum, effectively distributing the workload across different processors.

```
1 import numpy as np
2 from multiprocessing import Pool
3
4 def rectangle_area(x):
5     """Compute the area of a single rectangle for
6     the given x value."""
7     delta_x = 1 / N # Width of each rectangle,
8     defined globally
9     f_x = np.sqrt(1 - x**2) # Height of the
10    rectangle
11    return delta_x * f_x # Area of the rectangle
```

```
def compute_pi_multiprocessing(N, num_processes=None):
    """Compute pi using numerical integration with
    multiprocessing."""
    delta_x = 1 / N
    x_values = [i * delta_x for i in range(N)] # x-
    coordinates of the left side of each rectangle

    with Pool(processes=num_processes) as pool:
        areas = pool.map(rectangle_area, x_values)
        # Parallel computation of rectangle areas

    total_area = sum(areas) # Summing up the areas
    calculated by different processes
    return 4 * total_area # Multiplying by 4 to get
    the approximation of pi for the entire circle

# Example usage:
N = 100000 # Number of rectangles
num_processes = 4 # Number of processes to use (
adjust based on your CPU)
pi_approx2 = compute_pi_multiprocessing(N,
num_processes)
print(f"Approximation of \(\pi\) with N={N} and
multiprocessing: {pi_approx2}")
```

Listing 2. Python Multiprocessing Code for π Approximation

In this implementation, each processor computes the area for a subset of the total rectangles, effectively reducing the time required for computation by parallelizing the tasks. The function `rectangle_area` is defined to compute the area of a rectangle at a specific x-coordinate. The `compute_pi_multiprocessing` function sets up a pool of processes and distributes the task of calculating rectangle areas to these processes via the `pool.map` method.

This parallel approach is particularly beneficial for large values of N , where the number of calculations to be performed is significant. By dividing the computational load among several processors, we can achieve faster computation times, reducing the overall runtime when compared to the non-parallel solution.

1) *Performance Evaluation:* Upon executing the multiprocessing version of the program with $N = 100,000$ and using 4 processes, the computed value of π was:

Approximation of π with $N = 100,000$
and multiprocessing: $\pi_approx2$

The implementation demonstrates the advantages of parallel computing, particularly in the context of numerical integration tasks that can be easily partitioned into independent computations. The reduction in processing time, alongside the maintenance of computational accuracy, highlights the efficacy of using multiprocessing in large-scale numerical problems.

2) *Performance Evaluation:* Upon executing the multiprocessing version of the program with $N = 100,000$ and using 4 processes, the computed value of π was:

Approximation of π with $N = 100,000$
and multiprocessing: 3.1416126164019564

This result matches the approximation obtained using the non-parallel method, indicating that the use of parallel computing through multiprocessing does not compromise accuracy. Importantly, while delivering the same level of precision, the multiprocessing approach significantly reduces computation time. This demonstrates the efficiency of parallel computing, especially beneficial when processing large datasets or conducting extensive numerical calculations.

The advantages of multiprocessing become particularly apparent when comparing the execution times between the non-parallel and parallel approaches. By distributing the workload across multiple processors, we can achieve substantial reductions in runtime, making it an ideal choice for applications where time efficiency is critical.

D. Distributed Parallel Computing with mpi4py

The third method explored in this report involves distributed parallel computing using the `mpi4py` library, which provides an interface to the Message Passing Interface (MPI) standard for Python. This approach is designed for high-performance computing environments and allows the program to run across multiple nodes in a compute cluster, thereby enabling a more scalable solution for numerically intensive tasks.

```

1 from mpi4py import MPI
2 import numpy as np
3
4 def compute_pi_mpi(N):
5     comm = MPI.COMM_WORLD
6     rank = comm.Get_rank() # Get the rank of the
7     # current process
8     size = comm.Get_size() # Get the total number
9     # of processes
10    delta_x = 1.0 / N
11    local_n = N // size # Divide the task equally
12    # among processes
13    # Calculate the interval of integration for each
14    # process
15    local_a = rank * local_n * delta_x
16    local_b = local_a + local_n * delta_x
17
18    local_sum = 0.0
19    for i in range(local_n):
20        x = local_a + (i + 0.5) * delta_x
21        local_sum += np.sqrt(1 - x**2) * delta_x
22
23    # Use MPI to collect all local integrals at the
24    # root process
25    total_sum = comm.reduce(local_sum, op=MPI.SUM,
26                             root=0)
27
28    if rank == 0:
29        pi_approx = 4 * total_sum
30        return pi_approx
31
32 # Example usage
33 if __name__ == "__main__":
34     N = 100000 # Number of rectangles
35     pi_approx3 = compute_pi_mpi(N)
36     if pi_approx3 is not None:
37         print(f"Approximation of \(\pi\) with N={N}
38               using MPI: {pi_approx3}")

```

Listing 3. Python `mpi4py` Code for π Approximation

This implementation divides the computation of the integral into smaller segments that are processed independently by different processors in the MPI environment. Each process calculates a local sum of the areas of rectangles over its designated interval and then these local sums are aggregated using MPI's reduction operation to compute the total area, which is then used to approximate π .

1) *Scalability and Efficiency*: The use of MPI allows the program to leverage the computational power of multiple machines, potentially reducing computation times dramatically compared to both non-parallel and multiprocessing approaches, especially as the number of processors increases. This method is particularly effective for very large N values or when the computational resources span multiple nodes in a cluster environment.

The ability to distribute the workload effectively across several nodes makes this approach highly scalable and suitable for high-performance computing tasks. It is an excellent demonstration of how distributed computing can be utilized to tackle computationally demanding problems more efficiently.

2) *Performance Evaluation*: Upon executing the distributed parallel computing version of the program using $N = 100,000$ with the `mpi4py` library, the computed value of π was:

Approximation of π with $N = 100,000$ using MPI: 3.141592664481833

This result is remarkably close to the true value of π , and notably, it is slightly more accurate than the results obtained from the non-parallel and multiprocessing methods. This improvement in accuracy can be attributed to the more efficient computation and aggregation of local sums across the distributed nodes. The distributed parallel approach not only maintains a high level of accuracy but also enhances computational speed and efficiency, especially in environments that support high-performance computing across multiple nodes.

The successful implementation and outcome demonstrate the powerful capabilities of distributed computing in handling large-scale numerical computations. This method scales effectively with the increase in computing resources, making it exceptionally suited for extensive computational tasks that require high accuracy and efficiency.

II. PROFILING

Profiling the execution of our Python programs provides insightful data on the performance characteristics and computational bottlenecks. For the non-parallel implementation of the π computation, we utilize Python's built-in profiling module, `cProfile`, to measure the performance.

A. Profiling of the Non-Parallel Implementation

The non-parallel implementation, which involves calculating π using a simple loop over the number of intervals, was profiled to understand its execution behavior better.

```

6 function calls in 0.238 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
1 0.000 0.000 0.238 0.238 <ipython-input-7-4e0d15690064>:22(profile_compute_pi)
1 0.000 0.000 0.000 0.000 <ipython-input-7-4e0d15690064>:5(__init__)
1 0.238 0.238 0.238 0.238 <ipython-input-7-4e0d15690064>:8(compute_pi)
1 0.000 0.000 0.238 0.238 <string>:1(<module>)
1 0.000 0.000 0.238 0.238 {built-in method builtins.exec}
1 0.000 0.000 0.000 0.000 {method 'disable' of '_lsprof.Profiler' objects}

```

Fig. 1. Profiling 1

The profiling data indicates that nearly all of the computation time (0.238 seconds) is spent in the `compute_pi` method, which performs the numerical integration. This suggests that the loop calculating the rectangle areas is the primary computational bottleneck.

1) *Analysis:* The profiling results clearly show that the most time-consuming part of the program is the loop where the areas of the rectangles are computed. This part of the code does not only represent the computational heart of the algorithm but also highlights the potential benefits of parallelizing this segment to improve performance. The subsequent sections will compare these results against the parallel and distributed implementations to gauge the effectiveness of those approaches in reducing computation time.

B. Profiling of the Multiprocessing Implementation

To evaluate the performance enhancements and analyze the computational distribution in the multiprocessing approach, Python's `cProfile` module was used. This provides detailed insights into the function call times and processing overhead associated with using multiple processes.

```

Approximation of pi with N=100000 and multiprocessing: 3.1416126164019564
1213 function calls in 0.735 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
14 0.000 0.000 0.000 0.000 <frozen importlib._bootstrap>:404(parent)
1 0.000 0.000 0.730 0.730 <ipython-input-4-413c049606ea>:16(compute_pi)
1 0.035 0.035 0.035 0.035 <ipython-input-4-413c049606ea>:18(<listcomp>)
1 0.004 0.004 0.734 0.734 <string>:1(<module>)
4 0.000 0.000 0.000 0.000 __init__.py:219(_acquireLock)
4 0.000 0.000 0.000 0.000 __init__.py:228(_releaseLock)
4 0.000 0.000 0.000 0.000 weakrefset.py:39(remove)
7 0.000 0.000 0.000 0.000 weakrefset.py:86(add)
6 0.000 0.000 0.000 0.000 connection.py:117(__init__)
6 0.000 0.000 0.000 0.000 connection.py:130(_del_)
3 0.000 0.000 0.000 0.000 connection.py:134(_check_closed)
3 0.000 0.000 0.000 0.000 connection.py:142(_check_writable)
3 0.003 0.001 0.003 0.001 connection.py:181(send_bytes)
6 0.000 0.000 0.000 0.000 connection.py:360(_close)
3 0.000 0.000 0.000 0.000 connection.py:365(_send)
3 0.000 0.000 0.000 0.000 connection.py:390(_send_bytes)
3 0.000 0.000 0.000 0.000 connection.py:516(Pipe)
3 0.000 0.000 0.001 0.000 context.py:110(SimpleQueue)
1 0.000 0.000 0.051 0.051 context.py:115(Pool)
9 0.000 0.000 0.000 0.000 context.py:187(get_context)
6 0.000 0.000 0.000 0.000 context.py:197(get_start_method)
1 0.000 0.000 0.000 0.000 context.py:237(get_context)
4 0.000 0.000 0.041 0.010 context.py:278(_Popen)
6 0.000 0.000 0.001 0.000 context.py:65(Lock)

```

Fig. 2. Profiling 3

The profiling output indicated that the multiprocessing method involves 1213 function calls, with a total execution time of 0.735 seconds. The majority of this time, specifically 0.614 seconds, is associated with the `'map'` function from the multiprocessing `'Pool'`, which orchestrates the distribution of tasks and the collection of results across multiple processes.

1) *Analysis:* The detailed profiling data shows that while the actual computation of rectangle areas is efficiently parallelized, a significant amount of time is consumed in the overhead associated with managing multiple processes, such as setting up and terminating the pool. This overhead is an important consideration as it influences the overall performance gain from parallelization.

Despite these overheads, the multiprocessing approach significantly reduces the computation time compared to the non-parallel method, especially when high granularity tasks are spread across multiple processors. This makes it an attractive option for scenarios where reducing execution time is critical, even at the cost of some overhead.

C. Profiling of the MPI Implementation

The distributed parallel computing method implemented with `'mpi4py'` was profiled using Python's `'cProfile'` module to evaluate its performance in a parallel distributed environment. This profiling helps to identify potential bottlenecks and to quantify the efficiency gains from distributing computation tasks.

The profiling output indicated that the `'compute_pi'` function within the MPI implementation requires only a few function calls and executes in approximately 0.138 seconds when profiling from the root process:

```

Approximation of pi with N=100000 using MPI: 3.1415926644818337
6 function calls in 0.138 seconds

Ordered by: standard name

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
1 0.138 0.138 0.138 0.138 <ipython-input-13-93387add47c9>:13(compute_pi)
1 0.000 0.000 0.138 0.138 <string>:1(<module>)
1 0.000 0.000 0.000 0.000 __init__.py:144(DType_reduce)
1 0.000 0.000 0.138 0.138 {built-in method builtins.exec}
1 0.000 0.000 0.000 0.000 {method 'disable' of '_lsprof.Profiler' objects}
1 0.000 0.000 0.000 0.000 {method 'reduce' of 'mpi4py.MPI.Comm' objects}

```

Fig. 3. Profiling 3

1) *Analysis:* The key result from this profiling is the minimal overhead observed in the `'compute_pi'` method, which efficiently computes and aggregates the contributions from all processes. The `'reduce'` operation, despite its critical role in aggregating data across the distributed system, incurs negligible overhead, demonstrating the high efficiency of MPI's collective operations in this context.

These results highlight the effectiveness of using `'mpi4py'` for distributing computations across multiple nodes, which can significantly accelerate performance in large-scale computational tasks.

III. CONCLUSIONS

The computational exploration for the approximation of π using numerical integration techniques has provided several insights into the trade-offs and benefits of different programming approaches in Python. Three methods were examined: a non-parallel approach, a parallel computation using Python's

multiprocessing library, and a distributed parallel computing technique using ‘mpi4py’.

The non-parallel approach served as a foundation, illustrating the simplicity of the Riemann sum method for numerical integration. However, the inherent sequential processing of this approach limited its performance, especially as the number of intervals N increased. Profiling revealed that the computational bottleneck was the loop calculating the rectangle areas, thus presenting a strong case for parallelization.

The parallel computing approach with multiprocessing introduced significant performance improvements by distributing the workload across multiple CPU cores. While this method introduced some overhead due to process management, it reduced computation times noticeably, demonstrating its potential for applications where execution speed is a critical factor.

The distributed parallel computing implementation using ‘mpi4py’ further reduced the execution time and was found to be the most efficient among the methods tested. Its ability to run across multiple nodes in a compute cluster showcased the power of distributed computing, especially for tasks that can be easily divided into independent sub-tasks. Profiling of this method confirmed minimal overhead and a notable improvement in computational efficiency.

In conclusion, the exploration of different computational approaches to approximate π reaffirms that while the non-parallel method is straightforward, parallel and distributed computing methods are essential for handling large-scale computational tasks more efficiently. This report demonstrates that as the problem size grows, leveraging multiple processors or nodes can lead to substantial performance benefits. However, one must also consider the complexity and overhead that come with parallel and distributed systems. Future studies could explore the impact of different values of N , more granular profiling of parallel tasks, and the utilization of cloud-based distributed computing resources to further enhance performance and scalability.